



ECE408/CS483/CSE408 Spring 2023

## Applied Parallel Programming

# Lecture 9: Tiled Convolution Analysis

# Course Reminders


- Grades for Lab 1 are uploaded to Canvas soon
  - Check and let us know if they are missing
- Lab 4 out, it is due next week
- Midterm 1 is on Tuesday, March 7<sup>th</sup>
  - On-line, everybody will be taking it at the same time
    - Tuesday, March 7<sup>th</sup> 7:00pm-8:30pm US Central time
  - Includes materials from Lecture 1 through Lecture 9
- Project Milestone 1: Baseline CPU implementation is due Friday March 10<sup>th</sup>
  - Project details to be posted next week

# Objective

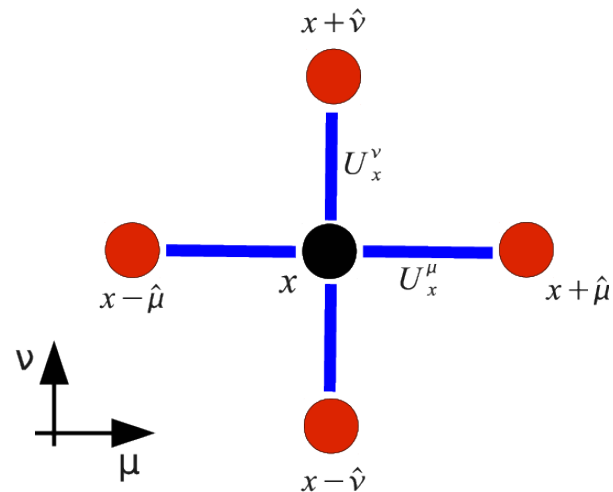
- To learn more about the analysis of tiled convolution/stencil algorithms

# Stencil Algorithms

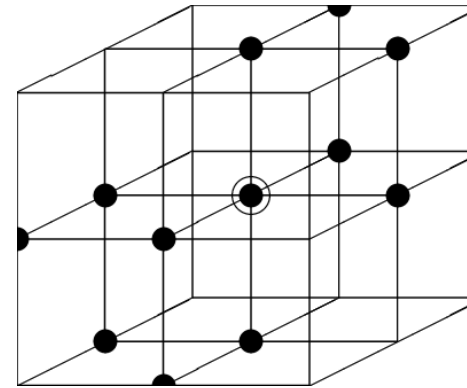
- Numerical data processing algorithms which update array elements according to some fixed pattern, called a *stencil*
  - Convolution is just one such example




2D convolutional kernel (9-point 2D stencil)



Nearest neighbor lattice Dirac operator (5-point 2D stencil)

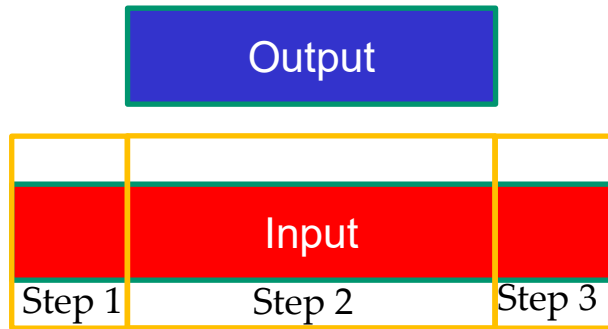


Finite difference stencil for 3D explicit time-marching (13-point 3D stencil)



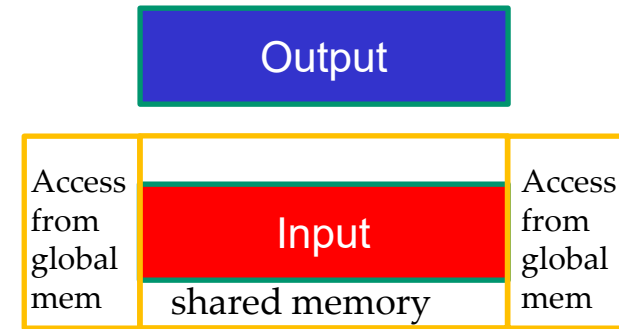
The Wilson-Dslash operator (4D stencil)

# Review: Three Tiling Strategies




Strategy 1

1. Block size covers **output** tile
2. Use multiple steps to load input tile



Strategy 3

1. Block size covers **output** tile
2. Load only “core” of input tile
3. Access halo cells from global memory




Strategy 2

1. Block size covers **input** tile
2. Load input tile in one step
3. Turn off some threads when calculating output

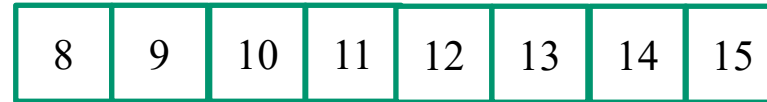
# A Small 1D Convolution Example

tile



MASK\_WIDTH is 5


P



TILE\_WIDTH is 8

- output and input tiles for block 1
- For MASK\_WIDTH of 5, each block loads  $8 + (5 - 1) = 12$  elements (**12 memory loads**)

# Each Output Uses MASK\_WIDTH Inputs



- P[8] uses N[6], N[7], N[8], N[9], N[10]
- P[9] uses N[7], N[8], N[9], N[10], N[11]
- ...
- P[15] uses N[13], N[14], N[15], N[16], N[17]

Total of **8 \* 5 values** from **tile** **used for the output.**

# A simple way to calculate tiling benefit

- $8+(5-1) = 12$  unique elements of input array N loaded
- $8*5$  global memory accesses potentially replaced by shared memory accesses
- This gives a bandwidth reduction of  $40/12=3.3$
- This is independent of the size of N




# In General, for 1D convolution kernels

- Load  $(\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)$  elements from global memory to shared memory
- Replace  $(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})$  global memory accesses with shared memory accesses
- This leads to bandwidth reduction of

$$(\text{TILE\_SIZE} * \text{MASK\_WIDTH}) / (\text{TILE\_SIZE} + \text{MASK\_WIDTH} - 1)$$

# Another Way to Look at Reuse



- tile[6] is used by P[8] (1×)
- tile[7] is used by P[8], P[9] (2×)
- tile[8] is used by P[8], P[9], P[10] (3×)
- tile[9] is used by P[8], P[9], P[10], P[11] (4×)
- tile[10] is used by P[8], P[9], P[10], P[11], P[12] (5×)
- ... (5×)
- tile[14] is used by P[12], P[13], P[14], P[15] (4×)
- tile[15] is used by P[13], P[14], P[15] (3×)
- tile[16] is used by P[14], P[15] (2×)
- tile[17] is used by P[15] (1×)


# Another Way to Look at Reuse

- Each access **tile** replaces an access to input N
- The total number of global memory accesses (to the  $(8+5-1)=12$  input elements) replaced by shared memory accesses is

$$\begin{aligned} & 1 + 2 + 3 + 4 + 5 * (8-5+1) + 4 + 3 + 2 + 1 \\ &= 10 + 20 + 10 \\ &= 40 \end{aligned}$$

- There are 12 elements of input N, so the average reduction is  
 **$40/12 = 3.3$**

# What about Boundary Tiles?



- P[0] uses N[0], N[1], N[2]
- P[1] uses N[0], N[1], N[2], N[3]
- P[2] uses N[0], N[1], N[2], N[3], N[4]

Less than  $8 * 5$  elements of N are used for the output tile

# Ghost elements change ratios

- For a boundary tile, we load  $\text{TILE\_WIDTH} + (\text{MASK\_WIDTH}-1)/2$  elements
  - 10 in our example of  $\text{TILE\_WIDTH}$  of 8 and  $\text{MASK\_WIDTH}$  of 5
- Computing boundary elements do not access global memory for ghost cells
  - Total accesses is  $6*5 + 4 + 3 = 37$  accesses (when computing the P elements)

The reduction is  **$37/10 = 3.7$**

# In General for 1D, internal tiles

- The total number of global memory accesses to the  $(\text{TILE\_WIDTH} + \text{Mask\_Width} - 1)$  elements of input N replaced by shared memory accesses is

$$\begin{aligned} &1 + 2 + \dots + \text{Mask\_Width} - 1 + \text{Mask\_Width} * (\text{TILE\_WIDTH} - \text{Mask\_Width} + 1) + \text{Mask\_Width} - 1 + \dots + 2 + 1 \\ &= ((\text{Mask\_Width} - 1) * \text{Mask\_Width}) / 2 + \text{Mask\_Width} * (\text{TILE\_WIDTH} - \text{Mask\_Width} + 1) + ((\text{Mask\_Width} - 1) * \text{Mask\_Width}) / 2 \end{aligned}$$

$$= (\text{Mask\_Width} - 1) * \text{Mask\_Width} + \text{Mask\_Width} * (\text{TILE\_WIDTH} - \text{Mask\_Width} + 1)$$

$$= \text{Mask\_Width} * \text{TILE\_WIDTH}$$

# Bandwidth Reduction for 1D


- The reduction is

$$(\text{TILE\_SIZE} * \text{MASK\_WIDTH}) / (\text{TILE\_SIZE} + \text{MASK\_WIDTH} - 1)$$

TILE_WIDTH	16	32	64	128	256
Reduction MASK_WIDTH = 5	4.0	4.4	4.7	4.9	4.9
Reduction MASK_WIDTH = 9	6.0	7.2	8.0	8.5	8.7

# Review: Parallelization of Tile Load

**MASK\_WIDTH is 5**





# Analysis for an 8×8 Output Tile, MASK\_WIDTH of 5

- Loading input tile requires  $(8+5-1)^2 = 144$  reads
- Calculation of each output requires  $5^2 = 25$  input elements
- $8 \times 8 \times 25 = 1,600$  global memory accesses for computing output tile are converted to shared memory accesses
- Bandwidth reduction of  $1,600/144 = 11.1\times$

# In General

- $(\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)^2$  elements need to be loaded from  $N$  into shared memory
- The calculation of each  $P$  element needs to access  $\text{MASK\_WIDTH}^2$  elements of  $N$
- $(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})^2$  global memory accesses converted into shared memory accesses
- Bandwidth reduction of  $(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})^2 / (\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)^2$

# Bandwidth Reduction for 2D Convolution Kernel

- The reduction is

$$(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})^2 / (\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)^2$$

TILE_WIDTH	8	16	32	64
Reduction MASK_WIDTH = 5	11.1	16	19.7	22.1
Reduction MASK_WIDTH= 9	20.3	36	51.8	64

# Ghost elements change ratios

- 2D calculation left for your enjoyment

# 2B/FLOP for Untiled Convolution

- How much global memory per FLOP is in untiled convolution?
- In untiled convolution,
  - each value from **N** (**4B** from global memory)
  - is multiplied by a value from **M** (**4B** from constant cache, **1 FLOP**),
  - then added to a running sum (**1 FLOP**)
- That gives **2B / FLOP**

# Full Use of Compute Requires 13.3× Reuse

- Recall our reuse discussion from matrix multiply:
  - **1,000 GFLOP/s** for GPU from ~2010, and
  - **150 GB/s** memory bandwidth.

- Dividing memory bandwidth by **2B/FLOP**,

$$\frac{150 \text{ GB/s}}{2 \text{ B/FLOP}} = \mathbf{75 \text{ GFLOP/s} = 7.50\% \text{ of peak.}}$$

- **Need** at least **100/7.50 = 13.3× reuse** to make full use of compute resources

# In 2020, Need 52.1× Reuse

- In 2020, the **GRID K520** offers
  - nearly **5,000 GFLOP/s**, but only
  - **192 GB/s** memory bandwidth

- Dividing memory bandwidth by **2B/FLOP**,

$$\frac{192 \text{ GB/s}}{2 \text{ B/FLOP}} = \mathbf{96 \text{ GFLOP/s} = 1.92\% \text{ of peak}}$$

- **Need** at least **100/1.92 = 52.1× reuse** to make full use of compute resources

# Need Really Big Mask to Balance Resources

- Let's make another table: **% of peak compute**
  - for **1D** tiled convolution,
  - with **TILE\_WIDTH 1024**

MASK_WIDTH	~2010 GPU with 1,000 GFLOP/s 150 GB/s	~2020 GPU with 5,000 GFLOP/s 192 GB/s
5	37%	9.6%
9	67%	17%
15	100%	28%
55	100%	100%



# Need Really Big Mask to Balance Resources

- And one more: **% of peak compute**
  - for **2D** tiled convolution,
  - with **TILE\_WIDTH 32×32**

MASK_WIDTH	~2010 GPU with 1,000 GFLOP/s 150 GB/s	~2020 GPU with 5,000 GFLOP/s 192 GB/s
3	60%	15%
5	100%	37%
7	100%	67%
9	100%	almost 100%

# Food for Thought


- Ratios are different for tiles on boundaries
- More importantly,
  - Each thread loads 4B to shared memory
  - 2,048 threads load only 8kB
  - Shared memory is usually 64kB or larger
  - What can one do with the rest?

**Improved approach left as homework**

(For example, can raise MW=7 from 67% to 81%)


# Problem Solving

- Q: Consider 2D tiled convolution with mask width of 5x5 and output tile width of 16x16 applied to an input image of size 32x32. Assume we use Strategy 1. The use of shared memory reduces the number of global memory accesses. What is the reduction in global memory accesses for thread block (0,0)?
  - A:
    - Count # of global memory accesses:  
 row 0:  $9 + 12 + 15 * 14$   
 row 1:  $12 + 16 + 20 * 14$   
 other 14 rows:  $15 + 20 + 25 * 14$
    - Count # of input elements to be used in the calculation:  $18 * 18$
- 
- The diagram illustrates a 32x32 input image (represented by a blue square). A 16x16 output tile (represented by a black square) is shown in the top-left corner. Within this output tile, a 5x5 mask (represented by a green square) is shown. The dimensions are labeled: the input image is 32x32, the output tile is 16x16, and the mask is 5x5. The distance from the top-left corner of the mask to the top-left corner of the output tile is 16+2 units. The distance from the bottom-left corner of the mask to the bottom-left corner of the output tile is 16+2 units. The distance from the top-right corner of the mask to the top-right corner of the output tile is 16+2 units. The distance from the bottom-right corner of the mask to the bottom-right corner of the output tile is 16+2 units.



# Problem Solving

- Q: For a tiled 2D convolution kernel with 30x30 output tiles and 3x3 mask (and thus 32x32 input tile), how many warps in each thread block have control divergence? (Assume strategy 2: block size covers input tile.)
- A:
  - thread block size is 32x32
  - All 32x32 threads participate in data load
  - What does each warp do?
    - Warp 0: load 0s and do not compute
    - Warps 1-30: 30 threads compute, and 2 threads do not
    - Warp 31: load 0s and do not compute
  - Thus, only 2 warps do not have control divergence



Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

**ANY MORE QUESTIONS?  
READ CHAPTER 7**